

```

        FILE: Number of write operations=0
        HDFS: Number of bytes read=184
        HDFS: Number of bytes written=69
        HDFS: Number of read operations=15
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=4
        HDFS: Number of bytes read erasure-coded=0
    Map-Reduce Framework
        Map input records=5
        Map output records=20
        Map output bytes=169
        Map output materialized bytes=215
        Input split bytes=102
        Combine input records=0
        Combine output records=0
        Reduce input groups=10
        Reduce shuffle bytes=215
        Reduce input records=20
        Reduce output records=10
        Spilled Records=40
        Shuffled Maps =1
        Failed Shuffles=0
        Merged Map outputs=1
        GC time elapsed (ms)=0
        Total committed heap usage (bytes)=1052770304
    Shuffle Errors
        BAD_ID=0
        CONNECTION=0
        IO_ERROR=0
        WRONG_LENGTH=0
        WRONG_MAP=0
        WRONG_REDUCE=0
    File Input Format Counters
        Bytes Read=92
    File Output Format Counters
        Bytes Written=69
0
hadoop@bnscecse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$ hadoop fs -ls /output/
Found 2 items
-rw-r--r--  1 hadoop supergroup          0 2025-04-15 14:46 /output/_SUCCESS
-rw-r--r--  1 hadoop supergroup        69 2025-04-15 14:46 /output/part-r-00000
hadoop@bnscecse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$ hadoop fs -cat /output/part-00000
cat: '/output/part-00000': No such file or directory
hadoop@bnscecse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$ hadoop fs -cat /output/part-r-00000
are      1
brother  1
family   1
hi        1
how       5
is        4
job       1
sister   1
you       1
your      4
hadoop@bnscecse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$

```

```

HDFS: Number of large read operations=0
HDFS: Number of write operations=3
HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Combine input records=0
  Combine output records=0
  Reduce input groups=10
  Reduce shuffle bytes=215
  Reduce input records=20
  Reduce output records=10
  Spilled Records=20
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=0
  Total committed heap usage (bytes)=526385152
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Output Format Counters
  Bytes Written=69
2025-04-29 15:41:08,074 INFO mapred.LocalJobRunner: Finishing task: attempt_local690373640_0001_r_000000_0
2025-04-29 15:41:08,074 INFO mapred.LocalJobRunner: reduce task executor complete.
2025-04-29 15:41:08,791 INFO mapreduce.Job: Job job_local690373640_0001 running in uber mode : false
2025-04-29 15:41:08,791 INFO mapreduce.Job: map 100% reduce 100%
2025-04-29 15:41:08,792 INFO mapreduce.Job: Job job_local690373640_0001 completed successfully
2025-04-29 15:41:08,796 INFO mapreduce.Job: Counters: 36
File System Counters
  FILE: Number of bytes read=8836
  FILE: Number of bytes written=1283959
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=184
  HDFS: Number of bytes written=69
  HDFS: Number of read operations=15
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=4
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Map input records=5
  Map output records=20
  Map output bytes=169
  Map output materialized bytes=215
  Input split bytes=102
  Combine input records=0
  Combine output records=0
  Reduce input groups=10
  Reduce shuffle bytes=215
  Reduce input records=20
  Reduce output records=10

```

```

copyFromLocal: /rpg/test.txt: File exists
hadoop@hncscse-HP-Eltite-Tower-600-G9-Desktop-PC: ~/Desktop$ hadoop jar /home/hduser/Desktop/WordCount170.jar
JAR does not exist or is not a normal file: /home/hduser/Desktop/WordCount170.jar
hadoop@hncscse-HP-Eltite-Tower-600-G9-Desktop-PC: ~/Desktop$ hadoop jar /home/hduser/Desktop/WordCount170.jar WCDriver input output
JAR does not exist or is not a normal file: /home/hduser/Desktop/WordCount170.jar
hadoop@hncscse-HP-Eltite-Tower-600-G9-Desktop-PC: ~/Desktop$ ls
ACFr0gCtAB80XZ9jzYb1bLxQ0_3-Mytt4Q3ozp_tew205dS9eh6tEqf_un0m1-d7mXh5B49Hx4Mg3zFkzS6AKBMTKHyxCYNBKeW7y0GfLl17d1B3CBKkdJx9GAsR23r3p2GQKyTrpg5UC5U.pdf
file1.txt
sample.txt
1735unit.jar
WordCount170.jar
WordCount.jar
hadoop@hncscse-HP-Eltite-Tower-600-G9-Desktop-PC: ~/Desktop$ hadoop jar WordCount170.jar WCDriver input output
2025-04-29 15:35:51.302 INFO Impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2025-04-29 15:35:51.341 INFO Impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2025-04-29 15:35:51.341 INFO Impl.MetricsSystemImpl: JobTracker metrics system started
2025-04-29 15:35:51.347 WARN Impl.MetricsSystemImpl: JobTracker metrics system already initialized!
2025-04-29 15:35:51.398 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2025-04-29 15:35:51.424 INFO mapreduce.JobSubmitter: Cleaning up the staging area file:/tmp/hadoop/mapred/staging/hadoop1877578914/staging/job_local1877578914_0001
Exception in thread "main" org.apache.hadoop.mapred.InvalidInputException: Input path does not exist: hdfs://localhost:9000/user/hadoop/input
    at org.apache.hadoop.mapred.FileInputFormat.singleThreadedListStatus(FileInputFormat.java:304)
    at org.apache.hadoop.mapred.FileInputFormat.listStatus(FileInputFormat.java:244)
    at org.apache.hadoop.mapred.FileInputFormat.getSplits(FileInputFormat.java:332)
    at org.apache.hadoop.mapreduce.JobSubmitter.writeOldSplits(JobSubmitter.java:336)
    at org.apache.hadoop.mapreduce.JobSubmitter.writeSplits(JobSubmitter.java:329)
    at org.apache.hadoop.mapreduce.JobSubmitter.submitJobInternal(JobSubmitter.java:200)
    at org.apache.hadoop.mapreduce.Job.submit(Job.java:1571)
    at org.apache.hadoop.mapreduce.Job$1.run(Job.java:1568)
    at java.base/java.security.AccessController.doPrivileged(Native Method)
    at java.base/java.security.auth.Subject.doAs(Subject.java:423)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1878)
    at org.apache.hadoop.mapred.JobClient$1.run(JobClient.java:576)
    at org.apache.hadoop.mapred.JobClient$1.run(JobClient.java:571)
    at java.base/java.security.AccessController.doPrivileged(Native Method)
    at java.base/java.security.auth.Subject.doAs(Subject.java:423)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1878)
    at org.apache.hadoop.mapred.JobClient.submitJobInternal(JobClient.java:571)
    at org.apache.hadoop.mapred.JobClient.submitJob(JobClient.java:562)
    at org.apache.hadoop.mapred.JobClient.runJob(JobClient.java:873)
    at WCDriver.run(WCDriver.java:30)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:81)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:95)
    at WCDriver.main(WCDriver.java:36)
    at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke(Native Method)
    at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
    at java.base/jdk.internal.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
    at java.base/java.lang.reflect.Method.invoke(Method.java:566)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:323)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:236)
Caused by: java.io.IOException: Input path does not exist: hdfs://localhost:9000/user/hadoop/input
    at org.apache.hadoop.mapred.FileInputFormat.singleThreadedListStatus(FileInputFormat.java:278)
    ... 29 more
hadoop@hncscse-HP-Eltite-Tower-600-G9-Desktop-PC: ~/Desktop$ bdfs dfs -ls input
Command 'bdfs' not found, did you mean:

```

```

    at org.apache.hadoop.util.RunJar.run(RunJar.java:323)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:236)
Caused by: java.io.IOException: Input path does not exist: hdfs://localhost:9000/user/hadoop/input
    at org.apache.hadoop.mapred.FileInputFormat.singleThreadedListStatus(FileInputFormat.java:278)
    ... 29 more
hadoop@hncscse-HP-Eltite-Tower-600-G9-Desktop-PC: ~/Desktop$ bdfs dfs -ls input
Command 'bdfs' not found, did you mean:
  command 'brfs' from deb node-brfs (2.0.2-2)
  command 'brfs' from deb brfs (2.24-1build1)
  command 'brs' from deb brs (2.2-1-1)
  command 'bzfs' from deb bzflag-server (2.4.22-1build1)
Try: sudo apt install <deb name>
hadoop@hncscse-HP-Eltite-Tower-600-G9-Desktop-PC: ~/Desktop$ hdfs dfs -ls input
ls: input: No such file or directory
hadoop@hncscse-HP-Eltite-Tower-600-G9-Desktop-PC: ~/Desktop$ hdfs dfs -mkdir input
mkdir: 'hdfs://localhost:9000/user/hadoop': No such file or directory
hadoop@hncscse-HP-Eltite-Tower-600-G9-Desktop-PC: ~/Desktop$ hdfs dfs -mkdir -p user/hadoop
hadoop@hncscse-HP-Eltite-Tower-600-G9-Desktop-PC: ~/Desktop$ hdfs dfs -put sample.txt input
hadoop@hncscse-HP-Eltite-Tower-600-G9-Desktop-PC: ~/Desktop$ hadoop jar WordCount170.jar WCDriver input output
2025-04-29 15:41:07.526 INFO Impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2025-04-29 15:41:07.562 INFO Impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2025-04-29 15:41:07.562 INFO Impl.MetricsSystemImpl: JobTracker metrics system started
2025-04-29 15:41:07.568 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2025-04-29 15:41:07.678 INFO mapred.FileInputFormat: Total input files to process : 1
2025-04-29 15:41:07.683 INFO mapreduce.JobSubmitter: number of splits:1
2025-04-29 15:41:07.737 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local690373640_0001
2025-04-29 15:41:07.737 INFO mapreduce.JobSubmitter: Executing with tokens: []
2025-04-29 15:41:07.787 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2025-04-29 15:41:07.788 INFO mapreduce.Job: Running job: job_local690373640_0001
2025-04-29 15:41:07.788 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2025-04-29 15:41:07.789 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapred.FileOutputCommitter
2025-04-29 15:41:07.791 INFO output.FileOutputCommitter: FileOutputCommitter Algorithm version is 2
2025-04-29 15:41:07.791 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2025-04-29 15:41:07.835 INFO mapred.LocalJobRunner: Waiting for map tasks
2025-04-29 15:41:07.836 INFO mapred.LocalJobRunner: Starting task: attempt_local690373640_0001_m_000000_0
2025-04-29 15:41:07.845 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2025-04-29 15:41:07.845 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2025-04-29 15:41:07.851 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
2025-04-29 15:41:07.859 INFO mapred.MapTask: Processing split: hdfs://localhost:9000/user/hadoop/input/sample.txt:0-92
2025-04-29 15:41:07.873 INFO mapred.MapTask: numReduceTasks: 1
2025-04-29 15:41:07.902 INFO mapred.MapTask: (EQUATOR) 0 kvl 26214396(104857584)
2025-04-29 15:41:07.902 INFO mapred.MapTask: mapreduce.task.sort.mb: 100
2025-04-29 15:41:07.902 INFO mapred.MapTask: soft limit at 83886080
2025-04-29 15:41:07.902 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
2025-04-29 15:41:07.902 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
2025-04-29 15:41:07.903 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
2025-04-29 15:41:07.946 INFO mapred.LineRecordReader: Found UTF-8 BOM and skipped it
2025-04-29 15:41:07.947 INFO mapred.LocalJobRunner:
2025-04-29 15:41:07.947 INFO mapred.MapTask: Starting flush of map output
2025-04-29 15:41:07.947 INFO mapred.MapTask: Spilling map output
2025-04-29 15:41:07.947 INFO mapred.MapTask: bufstart = 0; bufend = 169; bufvoid = 104857600
2025-04-29 15:41:07.947 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26214320(104857280); length = 77/6553600
2025-04-29 15:41:07.950 INFO mapred.MapTask: Finished spill 0

```



```

Starting namenagers
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~/hadoop/etc/hadoop$ cd -
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
localhost: namenode is running as process 15897. Stop it first and ensure /tmp/hadoop-hadoop-namenode.pid file is empty before retry.
Starting datanodes
localhost: datanode is running as process 16080. Stop it first and ensure /tmp/hadoop-hadoop-datanode.pid file is empty before retry.
Starting secondary namenodes [bmscscse-HP-Elite-Tower-600-G9-Desktop-PC]
bmscscse-HP-Elite-Tower-600-G9-Desktop-PC: secondarynamenode is running as process 16370. Stop it first and ensure /tmp/hadoop-hadoop-secondarynamenode.pid file is empty before retry.
Starting resourcemanager
resourcemanager is running as process 16661. Stop it first and ensure /tmp/hadoop-hadoop-resourcemanager.pid file is empty before retry.
Starting nodemanagers
localhost: nodemanager is running as process 16829. Stop it first and ensure /tmp/hadoop-hadoop-nodemanager.pid file is empty before retry.
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~$ jps
0256 org.eclipse.equinox.launcher_1.6.1000.v20250227-1734.jar
16080 DataNode
16370 SecondaryNameNode
16661 ResourceManager
15897 NameNode
16829 NodeManager
18399 Jps
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~$ ls
sktop      hadoop.txt      Public      wci.jar
.cocuments hs_err_pid4959.log  salpoorvika wc.jar
Downloads  hs_err_pid6247.log  share      WordCountClasses
eclipse-workspace hs_err_pid7049.log  snap      WordCountProject
hadoop     Music            Templates
hadoop-3.3.4.tar.gz Pictures         Videos
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~$ cd Desktop
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$ nano sample.txt
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$ hadoop fs -ls /
Found 4 items
drwxr-xr-x - hadoop supergroup          0 2024-05-14 15:15 /FF
drwxr-xr-x - hadoop supergroup          0 2024-05-14 14:58 /FFF
drwxr-xr-x - hadoop supergroup          0 2025-04-15 14:46 /output
drwxr-xr-x hadoop supergroup            0 2025-04-15 14:38 /rgs
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$ hadoop fs -mkdir /rgs
mkdir: '/rgs': File exists
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$ hadoop fs -copyFromLocal D:/sample.txt /rgs/test.txt
copyFromLocal: '/rgs/test.txt': File exists
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$ hadoop jar /home/hduser/Desktop/WordCount170.jar WCDriver input output
JAR does not exist or is not a normal file: /home/hduser/Desktop/WordCount170.jar
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$ hadoop jar /home/hduser/Desktop/WordCount170.jar WCDriver input output
JAR does not exist or is not a normal file: /home/hduser/Desktop/WordCount170.jar
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$ ls
ACFrOgC1t8BD0KZ9jzYb1xQB_3-Hytt4Q3oqpw_tewZOSdS9ehd1Eqf_un0n1-d7mXHSbB49Hx4MgZzFkzS6Ak8WTKHxycYNBKewr7yOGf1l7d1B3C8Kkdjx9GAsRZjr3pZGQKyTrpgSUCSU.pdf
file1.txt
sample.txt
rgs\test.txt.jar
WordCount170.jar
WordCount.jar
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~/Desktop$ hadoop jar WordCount170.jar WCDriver input output

```

```

hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~$ ls
sktop      hadoop.txt      Public      wci.jar
.cocuments hs_err_pid4959.log  salpoorvika wc.jar
Downloads  hs_err_pid6247.log  share      WordCountClasses
eclipse-workspace hs_err_pid7049.log  snap      WordCountProject
hadoop     Music            Templates
hadoop-3.3.4.tar.gz Pictures         Videos
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~$ cd hadoop
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~/hadoop$ ls
bin      libexec  LICENSE.txt  NOTICE-binary  shbin  WMC.txt
etc      LICENSE-binary  logs        NOTICE.txt      share
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~/hadoop$ cd etc
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~/hadoop/etc$ ls
hadoop
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~/hadoop/etc$ cd hadoop
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~/hadoop/etc/hadoop$ ls
capacity-scheduler.xml      kms-log4j.properties
configuration.xml           kms-site.xml
container-executor.cfg      log4j.properties
core-site.xml               mapred-env.cmd
hadoop-env.cmd              mapred-env.sh
hadoop-env.sh               mapred-queues.xml.template
hadoop-metrics2.properties  mapred-site.xml
hadoop-policy.xml           shellprofile.d
hadoop-user-functions.sh.example ssl-client.xml.example
fs-rbf-site.xml             ssl-server.xml.example
hdfs-site.xml               user_rc_policies.xml.template
https-env.sh                workers
https-log4j.properties     yarn-env.cmd
https-site.xml              yarn-env.sh
kms-acls.xml                yarnservice-log4j.properties
kms-env.sh                  yarn-site.xml
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~/hadoop/etc/hadoop$ nano mapred-site.xml
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~/hadoop/etc/hadoop$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [bmscscse-HP-Elite-Tower-600-G9-Desktop-PC]
Starting resourcemanager
Starting nodemanagers
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~/hadoop/etc/hadoop$ cd -
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC:~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
localhost: namenode is running as process 15897. Stop it first and ensure /tmp/hadoop-hadoop-namenode.pid file is empty before retry.
Starting datanodes
localhost: datanode is running as process 16080. Stop it first and ensure /tmp/hadoop-hadoop-datanode.pid file is empty before retry.
Starting secondary namenodes [bmscscse-HP-Elite-Tower-600-G9-Desktop-PC]
bmscscse-HP-Elite-Tower-600-G9-Desktop-PC: secondarynamenode is running as process 16370. Stop it first and ensure /tmp/hadoop-hadoop-secondarynamenode.pid file is empty before retry.
Starting resourcemanager

```